

Appendix B

Model Specifications for an ABM of Brokerage in Matching Markets

This document specifies an agent-based model of brokered matching.

The model is a discrete-time simulation of a matching market with *agents* and a *broker*. Agents seeking pairwise matches either search on their own or outsource to the broker. All economic quantities (match output, fees, surplus) are in the same payoff units. All agents use heuristic decision rules. No agent holds beliefs about other agents' strategies.

The model specification covers agents (§1), the matching problem (§2), how agents learn to predict expected match output (§3), match economics (§4), network structure and agent turnover (§5), search and broker access (§6), the outsourcing decision (§7), and period measures (§8). Baseline parameter values and reproducibility are documented in Part II (§9).

Part I. Model Specification

1. Agents and Time

The model runs for T discrete periods (default 500). It has N agents (default 1000), a single broker, and an undirected network G over the agents and broker.

Each agent i is characterized by:

- **Type** $\mathbf{x}_i \in \mathbb{R}^d$: a fixed vector of observable characteristics, drawn as specified below and used by the matching function (§2).
- **Current-period matches** M_i^t : active bilateral relationships involving i in period t . Each counterparty can appear at most once in M_i^t ; incoming accepted offers are not capped.
- **Active demand bound** K : the maximum number of active demands an agent can initiate in a period (default $K = 5$; §6).
- **Experience history** $\mathcal{H}_i^t = \{(\mathbf{x}_j, q_{ij})\}$: (other party's type, realized match output) pairs from all matches involving i (§3a).
- **Satisfaction indices** $s_{i,c}^t$: one scalar per search channel $c \in \{\text{self}, \text{broker}\}$, used in the outsourcing decision (§7).

Agent Types

Type vectors are multidimensional but concentrated near a smooth one-dimensional curve on the unit sphere in $\mathbb{R}^d$. Let $u \in [0, 1]$ index position along the curve:

$$\gamma(u) = \frac{\tilde{\gamma}(u)}{\|\tilde{\gamma}(u)\|}, \quad \tilde{\gamma}_\ell(u) = \begin{cases} \sin(2\pi\nu_\ell u + \psi_\ell) & \ell = 1, \dots, d_\gamma \\ 0 & \ell = d_\gamma + 1, \dots, d. \end{cases}$$

The frequencies $\nu_\ell \sim \text{Unif}\{1, 2, 3, 4, 5\}$ and phases $\psi_\ell \sim \text{Unif}[0, 2\pi)$ are drawn once per simulation. The parameter $d_\gamma \leq d$ is the number of coordinates in which the curve varies.

Each agent is drawn at a random position $u_i \sim \text{Unif}[0, 1]$ on the curve, then perturbed:

$$\mathbf{x}_i = \frac{\gamma(u_i) + \boldsymbol{\epsilon}_i}{\|\gamma(u_i) + \boldsymbol{\epsilon}_i\|}, \quad \boldsymbol{\epsilon}_i \sim \mathcal{N}\left(\mathbf{0}, \frac{\sigma_x^2}{d} \mathbf{I}_d\right).$$

Noise is applied in all d dimensions, then the vector is re-normalized. The scale $\sigma_x/\sqrt{d}$ keeps the expected displacement from the curve approximately σ_x for any d .

Thus the curve remains one-dimensional for every d_γ . Increasing d_γ changes how many coordinates contain systematic curve variation; the remaining coordinates contain only perturbation noise.

Broker

The broker is a permanent node in G and has no type vector. Its state is:

- **Experience history** $\mathcal{H}_b^t = \{(\mathbf{x}_i, \mathbf{x}_j, q_{ij})\}$: (party 1 type, party 2 type, realized match output) triples from mediated matches (§3c). The party labels have no directional meaning because the broker's pair features are symmetric.
- **Roster** Roster^t : the broker's standing access base (§6b).
- **Current clients** D^t : agents who outsource to the broker in period t (§6b).
- **Reputation** rep^t : the average broker satisfaction of current clients (§7).

2. The Matching Problem

Each potential pair's expected match output combines two sources: the agents' average general quality and a match-specific complementarity between their types. The parameter ρ moves the matching problem from pure complementarity to pure general quality. The parameter δ makes complementarity harder to predict by increasing the difference between otherwise similar pairs on opposite sides of a boundary in type-pair space. Expected and realized match output are symmetric in the two agents.

2a. Match Output

Let q_{ij} be the realized per-period output of a match between agents i and j , measured in payoff units:

$$q_{ij} = Q + f(\mathbf{x}_i, \mathbf{x}_j) + \varepsilon_{ij}, \quad \varepsilon_{ij} \sim \mathcal{N}(0, \sigma_\varepsilon^2).$$

The conditional expected output is $E[q_{ij} \mid \mathbf{x}_i, \mathbf{x}_j] = Q + f(\mathbf{x}_i, \mathbf{x}_j)$. The systematic component f is fully determined by the two types and by matching-function objects fixed at initialization. The

independent, mean-zero shock ε_{ij} creates match-specific variation that neither the agents nor the broker can predict.

The function $f : \mathbb{R}^d \times \mathbb{R}^d \rightarrow \mathbb{R}$ is unknown to all agents:

$$f(\mathbf{x}_i, \mathbf{x}_j) = \frac{1}{s_{\text{mix}}(\rho, \delta)} \left[\sqrt{\rho} \tilde{U}_{ij} + \sqrt{1 - \rho} \tilde{W}_{ij}(\delta) \right].$$

The tildes indicate that general quality (§2b) and match-specific complementarity (§2c) are centered and scaled using fixed moments. These moments are estimated once from independently generated matching environments. Within each environment, they are calculated across all unordered pairs of distinct agents; they are then averaged across environments and held fixed in every simulation. Individual markets are not rescaled, so they retain ordinary between-market variation.

Let $\kappa(\delta)$ denote the fixed covariance between general quality and complementarity after each component is standardized. The common adjustment

$$s_{\text{mix}}^2(\rho, \delta) = 1 + 2\sqrt{\rho(1 - \rho)} \kappa(\delta)$$

accounts for this small covariance. Averaged across the independently generated environments, each standardized component and the resulting systematic component f have within-market variance one for every ρ and δ . Because Q is constant, expected match output has the same variance. This defines the payoff unit. In these units, $Q = 2.84$ and $\sigma_\varepsilon = 0.28$.

Because the two components have the same marginal variance on this scale, $\rho \in [0, 1]$ is general quality's share of their summed marginal variances, while $1 - \rho$ is complementarity's share. Thus $\rho = 0$ gives pure complementarity and $\rho = 1$ gives pure general quality. At $\rho = 1$, δ has no effect.

2b. Agent General Quality

$\mathbf{c}$ is an ideal type vector. An agent's general-quality score, $\mathbf{x}_i^\top \mathbf{c}$, measures its alignment with that ideal. Pair-level general quality is the average of the two agents' scores:

$$U_{ij} = \frac{1}{2}(\mathbf{x}_i^\top \mathbf{c} + \mathbf{x}_j^\top \mathbf{c}), \quad \tilde{U}_{ij} = \frac{U_{ij} - \mu_U}{\sigma_U}.$$

The ideal type $\mathbf{c} \in \mathbb{R}^d$ is drawn at initialization as a perturbation of a random point on the agent type curve:

$$u_c \sim \text{Unif}[0, 1], \quad \boldsymbol{\epsilon}_c \sim \mathcal{N}\left(\mathbf{0}, \frac{\sigma_x^2}{d} \mathbf{I}_d\right), \quad \mathbf{c} = \frac{\boldsymbol{\gamma}(u_c) + \boldsymbol{\epsilon}_c}{\|\boldsymbol{\gamma}(u_c) + \boldsymbol{\epsilon}_c\|_2}.$$

The normalization places $\mathbf{c}$ on the same unit sphere as agent types. The perturbation therefore changes the ideal direction without changing its magnitude.

2c. Match-Specific Complementarity

Define base complementarity, signed complementarity, and difficulty-adjusted complementarity as

$$V_{ij} = \mathbf{x}_i^\top \mathbf{A} \mathbf{x}_j, \quad H_{ij} = \text{sign}(\mathbf{x}_i^\top \mathbf{B} \mathbf{x}_j) V_{ij}, \quad W_{ij}(\delta) = V_{ij} + \delta H_{ij}.$$

The normalized complementarity used in the matching function is

$$\widetilde{W}_{ij}(\delta) = \frac{W_{ij}(\delta) - \mu_W(\delta)}{\sigma_W(\delta)}.$$

$\mathbf{A}$ is a symmetric positive definite matrix that gives a smooth bilinear complementarity geometry. Because V_{ij} depends jointly on both types, it cannot in general be represented as the sum of separate agent scores. A second symmetric matrix, $\mathbf{B}$, assigns pairs to two regions of type-pair space. Both matrices are fixed for the duration of the simulation.

Base complementarity. Draw $\mathbf{M}_A \in \mathbb{R}^{d \times d}$ with iid $\mathcal{N}(0, 1)$ entries and set

$$\mathbf{A} = \mathbf{M}_A^\top \mathbf{M}_A \cdot \frac{d}{\text{tr}(\mathbf{M}_A^\top \mathbf{M}_A)}.$$

Thus $\mathbf{A}$ is symmetric positive definite and $\text{tr}(\mathbf{A}) = d$.

The trace normalization prevents arbitrary scale differences in $\mathbf{A}$ across realizations. The fixed pair-level normalization described in §2a then places complementarity on the common scale used in the matching function.

Difficulty parameter. Equivalently, write $W_{ij}(\delta) = g(\mathbf{x}_i, \mathbf{x}_j)V_{ij}$, where a second symmetric matrix $\mathbf{B} \in \mathbb{R}^{d \times d}$ defines

$$g(\mathbf{x}_i, \mathbf{x}_j) = 1 + \delta \cdot \text{sign}(\mathbf{x}_i^\top \mathbf{B} \mathbf{x}_j)$$

where $\delta \in [0, 1]$ (default 0.5).

The multiplier is $1 + \delta$ when $\mathbf{x}_i^\top \mathbf{B} \mathbf{x}_j > 0$, $1 - \delta$ when it is negative, and 1 when it is zero. At $\delta = 0$, the boundary between the two regions has no effect. For $\delta > 0$, pairs on opposite sides of the boundary can have sharply different expected output even when their types are otherwise similar. Increasing δ increases this difference, making complementarity harder to learn. Normalizing $W_{ij}(\delta)$ separately for each δ removes the accompanying change in dispersion. Thus δ changes the shape and difficulty of complementarity, not its variance.

Boundary construction. Matrix $\mathbf{B}$ is made orthogonal to $\mathbf{A}$ over the realized type distribution. This ensures that δ introduces a distinct boundary rather than merely amplifying the base complementarity pattern. Let

$$\Sigma_x = \frac{1}{N} \sum_{i=1}^N \mathbf{x}_i \mathbf{x}_i^\top$$

be the uncentered second-moment matrix of the realized type vectors.

Draw $\mathbf{Z} \in \mathbb{R}^{d \times d}$ with iid $\mathcal{N}(0, 1)$ entries and form a symmetric zero-trace matrix:

$$\tilde{\mathbf{B}} = \frac{\mathbf{Z} + \mathbf{Z}^\top}{2}, \quad \mathbf{B}_0 = \tilde{\mathbf{B}} - \frac{\text{tr}(\tilde{\mathbf{B}})}{d} \mathbf{I}_d.$$

Define the weighted inner product

$$\langle \mathbf{M}, \mathbf{N} \rangle_{\Sigma_x} = \text{tr}(\Sigma_x \mathbf{M} \Sigma_x \mathbf{N}),$$

then remove the projection of $\mathbf{B}_0$ onto $\mathbf{A}$ and normalize:

$$\mathbf{B}_{\text{raw}} = \mathbf{B}_0 - \frac{\text{tr}(\Sigma_x \mathbf{B}_0 \Sigma_x \mathbf{A})}{\text{tr}(\Sigma_x \mathbf{A} \Sigma_x \mathbf{A})} \mathbf{A}, \quad \mathbf{B} = \frac{\mathbf{B}_{\text{raw}}}{\|\mathbf{B}_{\text{raw}}\|_F}.$$

This gives $\text{tr}(\Sigma_x \mathbf{B} \Sigma_x \mathbf{A}) = 0$ and $\|\mathbf{B}\|_F = 1$. Weighting by Σ_x removes overlap between the two bilinear scores where agents actually occur, rather than in unused regions of $\mathbb{R}^d$.

Symmetry of $\mathbf{A}$ and $\mathbf{B}$ makes $f(\mathbf{x}_i, \mathbf{x}_j) = f(\mathbf{x}_j, \mathbf{x}_i)$.

3. Learning

Before searching for matches, agents and the broker may update their prediction models on their accumulated histories. Updated models are used to score candidates (§6). The broker is eligible to refit every period. Agent i is eligible when $i \bmod 2 = t \bmod 2$, so each agent is eligible every other period. An eligible model refits only when its history is nonempty and it has received new observations since its previous fitting call. A parameter switch selects the prediction model. Neural-network learning is the default. Ridge regression is the alternate experiment.

3a. Prediction Models and Fitting

In the default model, both agents and the broker use a fully-connected network with one hidden layer, ReLU activations, and a single linear output:

$$\hat{q}(\mathbf{z}) = \mathbf{w}_2^\top \text{ReLU}(\mathbf{W}_1 \mathbf{z} + \mathbf{b}_1) + b_2$$

where $\mathbf{z}$ is the input feature vector and the remaining symbols are network weights and biases. Agents receive raw partner type vectors, the broker receives pairs of them.

Training rule. Network weights minimize MSE with the Adam optimizer on the current training window. Agents and the broker may have separate learning rates, $\eta_{\text{lr},a}$ and $\eta_{\text{lr},b}$. No explicit weight regularization is applied. At $t = 0$, each network sets $\mathbf{W}_1$ using He initialization, $\mathbf{b}_1 = \mathbf{0}$, $\mathbf{w}_2 = \mathbf{0}$, and $b_2 = Q$ (§2a). An untrained predictor therefore returns exactly Q for every input. Networks with seed observations then take $E_{\text{init},c}$ gradient steps on their seed histories, where $c \in \{a, b\}$.

Alternate Ridge rule. In the Ridge experiment, agents and the broker instead use

$$\hat{q}_i(\mathbf{x}_j) = \alpha_i + \beta_i^\top \mathbf{x}_j, \quad \hat{q}_b(\mathbf{x}_i, \mathbf{x}_j) = \alpha_b + \gamma_b^\top \psi(\mathbf{x}_i, \mathbf{x}_j).$$

Agent and broker fits minimize mean squared error plus λ_a and λ_b , respectively, times the squared Euclidean norm of the slope coefficients. Intercepts are not penalized. Raw type and pair features are used without sample standardization. Before its first fit, a Ridge predictor has zero slopes and intercept Q . The reporting simulations use $\lambda_a = 0.003$ and $\lambda_b = 0.001$.

Subsequent NN updates warm-start from the previous weights and persistent Adam moment state. The number of gradient steps adapts to the ratio of new observations to total history:

$$E_{c,t} = \max\left(E_{\text{steps},c}, \left\lceil E_{\text{init},c} \cdot \frac{n_{\text{new}}}{n_{\text{total}}} \right\rceil\right), \quad c \in \{a, b\}.$$

where n_{new} is the number of observations added since that model’s previous fitting call, $n_{\text{total}} = |\mathcal{H}^t|$ is full history size, and $E_{\text{steps},c}$ is the learner-specific minimum per-call step count. For an agent, n_{new} can span multiple periods because agents fit on an alternating schedule. Agents and the broker use the same functional rule but may have different parameter values. Ridge is solved directly at each fitting call. The reporting NN values are $\eta_{\text{lr},a} = 0.003$, $\eta_{\text{lr},b} = 0.03$, $E_{\text{init},a} = E_{\text{init},b} = 100$, and $E_{\text{steps},a} = E_{\text{steps},b} = 50$.

Initialization is period 0 for each learner. Its seed-history count is stored as the first period boundary. Each later boundary stores cumulative history at the end of a simulation period. A fitting call uses observations from the most recent $W_p = 40$ completed simulation periods. Before 40 simulation periods are available, the window also includes period 0. Once 40 are available, the period-0 seed history is excluded. If the selected history block exceeds $M_{\text{train}} = 2000$ observations, fitting deterministically keeps 2000 spaced observations across that block, including its oldest and newest observations. If the cap is one, it keeps the newest observation. The cap changes only the fitting sample, not the stored history or the period window. NN and Ridge use the same window, cap, fitting cadence, and histories.

3b. Agent i ’s Model

History. $\mathcal{H}_i^t = \{(\mathbf{x}_j, q_{ij})\}_{m=1}^{n_i}$ records the other party’s type and realized match output from every match involving i .

Input. The agent’s network takes the partner’s type as input: $\mathbf{z} = \mathbf{x}_j$.

The hidden width is $h_A = 2d$, so the baseline uses 16 hidden units.

3c. Broker’s Model

History. $\mathcal{H}_b^t = \{(\mathbf{x}_i, \mathbf{x}_j, q_{ij})\}_{m=1}^{n_b}$ records both parties’ types and realized match output from broker-mediated matches.

Input. The broker’s network takes a symmetric pair feature vector:

$$\psi(\mathbf{x}_i, \mathbf{x}_j) = \left[\mathbf{x}_i + \mathbf{x}_j; \text{vech}\left(\frac{\mathbf{x}_i \mathbf{x}_j^\top + \mathbf{x}_j \mathbf{x}_i^\top}{2}\right) \right].$$

The half-vectorization vech keeps the lower-triangular entries of the symmetric matrix, so the broker input dimension is $d_B = d + d(d+1)/2$, equal to 44 at $d = 8$.

The broker hidden width is $h_B = 8d$, so the baseline uses 64 hidden units.

Because $\psi(\mathbf{x}_i, \mathbf{x}_j) = \psi(\mathbf{x}_j, \mathbf{x}_i)$, each observed unordered pair contributes one training row.

The broker learns only from mediated matches. It does not observe self-search outcomes.

3d. Ridge Broker Ablations

Three separate Ridge-only ablations change the broker while leaving every agent’s local Ridge model unchanged. **Size-matched pooled history** uniformly samples the broker’s windowed and capped fitting data without replacement, using the smaller of its available sample size and the nearest-integer median sample size among agents with nonempty histories, with half-integers rounded up. **Single-principal history** randomly retains one endpoint when each broker observation is recorded,

fits $g_b(\mathbf{x}) = \alpha_b + \gamma_b^\top \mathbf{x}$, and scores a pair as $g_b(\mathbf{x}_i) + g_b(\mathbf{x}_j) - \bar{q}_b$, where $\bar{q}_b$ is the current fitting-sample output mean and defaults to Q before fitting. **Additive Ridge** uses $\mathbf{x}_i + \mathbf{x}_j$ without bilinear features. All use the calibrated agent and broker penalties, period window, observation cap, and fitting cadence.

4. Match Economics

Matches use a common reservation threshold and channel-specific cost accounting.

At initialization, 10,000 random agent pairs define the calibration mean $\bar{q}_{\text{cal}} = E[q]$. This calibration quantity scales r , ϕ , and c_s .

4a. Outside Options

All agents share a reservation threshold r , calibrated at initialization:

$$r = f_r \cdot \bar{q}_{\text{cal}}$$

where $\bar{q}_{\text{cal}}$ is the calibration mean and f_r is the reservation coefficient (default 0.60). Since ϕ and c_s do not depend on r (§4c), f_r may take any nonnegative value, including $f_r \geq 1$, and is swept as an analysis axis.

4b. Participation Constraints

Demanders emit offers only above the reservation threshold. Self-search uses the demander evaluation $\hat{q}_i(\mathbf{x}_j) > r$. Brokered search uses the broker pair prediction $\hat{q}_b(\{i, j\}) > r$. A one-sided offer is accepted only if the receiver’s evaluation exceeds r ; reciprocal offers are accepted automatically (§6d).

4c. Search Costs

The self-search cost and broker fee share one calibrated level:

$$\phi = c_s = \lambda_c \cdot \bar{q}_{\text{cal}}.$$

λ_c is the shared cost rate (default 0.05). The two symbols differ in payoff accounting:

- **Self-search** costs c_s on each requested self-search position, whether or not it is filled.
- **Brokerage** charges ϕ only on successful brokered placements.

5. Network Structure and Turnover

Agents interact through a single undirected network G that determines search opportunities and structural position.

5a. Network Initialization

G is initialized as a Watts-Strogatz small-world graph on the N agents (Watts & Strogatz, 1998), with ring degree $k_G = 6$, rewiring probability $p_{\text{rewire}} = 0.1$, and random agent order on the ring.

The broker is added as a permanent node with no type vector. Broker-node edges are synchronized to standing roster members, current-period broker clients, and agents currently engaged in broker-channel matches (§6b).

5b. Match Tie Formation

Each realized match adds an undirected i - j edge if one does not already exist. Ties persist until either node exits.

5c. Agent Turnover

Agents exit independently each period with probability η_{exit} (default 0.02). An exiting agent is removed everywhere. Its node slot is reused by an entrant with a fresh type drawn from the same type process (§1). The entrant receives $\lfloor k_G/2 \rfloor$ edges to existing agents sampled with probability proportional to $\exp(-\|\mathbf{x}_{i'} - \mathbf{x}_j\|^2)$.

6. Search

At the start of each period, each agent draws active relationship demand

$$d_i \sim \text{Binomial}(K, p_{\text{demand}}),$$

where K is the maximum number of active demands the agent can initiate in the period and p_{demand} defaults to 0.50. If $d_i > 0$, the agent chooses **one channel for the batch** (§7): self-search or broker.

6a. Self-Search

Agent i 's self-search candidate pool has two components:

Neighbors. For a direct network neighbor j with prior pair-specific experience, agent i uses $\bar{q}_{ij} = \frac{1}{n_{ij}} \sum q_{ij}^{(m)}$, where n_{ij} is the number of times i and j have matched. If the graph edge exists but $n_{ij} = 0$, the agent instead uses $\hat{q}_i(\mathbf{x}_j)$. Initial histories seed one observation for every initial graph edge. Unobserved neighbor edges arise after turnover because entrants attach structurally before matching.

Strangers. Each self-searching demander i independently samples a pool U_i^t of $\min(n_{\text{strangers}}, N)$ agents uniformly without replacement, where $n_{\text{strangers}} = 10$ (default). The demander evaluates pool members that are not current neighbors using $\hat{q}_i(\mathbf{x}_j)$ (§3b).

Agent i orders feasible self-search candidates by these evaluations and emits up to d_i directed offers to candidates above r . Equal evaluations are ordered uniformly at random, so an agent whose predictions are constant produces a random ranking rather than an undefined ordering.

6b. Broker Roster and Access

The broker maintains a **roster** of agents it can propose as counterparties when mediating matches. In each period, it also observes the current-client set D^t of agents who outsourced. Broker access for the period is:

$$\mathcal{A}_b^t = \text{Roster}^t \cup D^t.$$

Current clients expand broker access for the period but do not become standing roster members.

Initialization. The roster is seeded with a target size

$$R^* = \lceil \alpha_R N \rceil,$$

where the broker roster share α_R defaults to 0.20. The model draws R^* agents uniformly at random (default 200 at $N = 1000$). The broker’s history is seeded from existing roster-roster ties in the initial graph, capped at 100 observed ties.

Roster refresh. At the start of each period, each roster member independently exits with probability p_{roster} (default 0.02). If $\widetilde{\text{Roster}}^t$ is the post-churn roster,

$$\widetilde{\text{Roster}}^t = \{i \in \text{Roster}^{t-1} : u_i^t > p_{\text{roster}}\}, \quad u_i^t \stackrel{iid}{\sim} \text{Unif}[0, 1],$$

then the broker samples without replacement from $\{1, \dots, N\} \setminus \widetilde{\text{Roster}}^t$ until $|\text{Roster}^t| = R^*$, or until the population is exhausted. Broker-node edges are synchronized to the standing roster, current clients, and current broker-mediated matches (§5a).

6c. Broker-Mediated Search

For every unordered pair $\{i, j\}$ such that one side is a broker-client demander and the other side is in $\mathcal{A}_b^t$, the broker computes predicted match output. For each broker-client demander i , it ranks accessible counterparties above r , breaking equal predictions uniformly at random, and keeps the highest-ranked candidates up to i ’s residual active broker demand. Directed broker offers are then processed in a global score ranking with a new, independent random ordering within score ties. If both sides are broker-client demanders with remaining active demand, the same unordered pair can generate reciprocal broker offers.

The broker-service experiment separates access from assessment. The default **full-service broker** uses the broker access set $\mathcal{A}_b^t$ and the broker’s pair predictions. The **assessment-only broker** applies the broker’s pair predictions to the same candidate set an agent could reach through self-search: its network neighbors and an independently drawn stranger pool. The **access-only broker** presents candidates from $\mathcal{A}_b^t$, but the client orders them using its own assessments: its pair-specific mean for an experienced neighbor and its prediction model otherwise. In every mode, candidates below the score source’s reservation threshold are excluded.

These restrictions change only how brokered candidates are obtained and ordered. The broker continues to observe mediated outcomes and refit its model in all modes, although its predictions do not determine access-only offers. Offer acceptance, match formation, fees, satisfaction, reputation, learning histories, and network updating are unchanged. “Assessment-only” means that the broker adds no candidate access beyond ordinary self-search. It does not imply that every resulting match

joins existing neighbors, because ordinary self-search includes strangers. Likewise, access-only service can select a current neighbor when that neighbor belongs to the broker access set.

6d. Offer Acceptance and Match Formation

All self-search and broker offers enter one shared market, grouped by unordered pair.

1. If pair $\{i, j\}$ contains reciprocal offers $i \rightarrow j$ and $j \rightarrow i$, the relationship is accepted automatically.
2. If pair $\{i, j\}$ contains one directed offer $i \rightarrow j$, receiver j evaluates i using $\bar{q}_{ji}$ for known partners and $\hat{q}_j(\mathbf{x}_i)$ otherwise. The offer is accepted iff this value exceeds r .
3. Each accepted unordered pair forms at most one current-period relationship.

At market finalization, for each accepted match: 1. Realized output is drawn: $q_{ij} = Q + f(\mathbf{x}_i, \mathbf{x}_j) + \varepsilon_{ij}$. 2. Both parties add the observation to their histories: i adds $(\mathbf{x}_j, q_{ij})$ to $\mathcal{H}_i$; j adds $(\mathbf{x}_i, q_{ij})$ to $\mathcal{H}_j$. 3. If either directed offer used the broker channel, the broker adds $(\mathbf{x}_i, \mathbf{x}_j, q_{ij})$ to $\mathcal{H}_b$ once. 4. An edge is added between i and j in G (if not already present).

A demander can have at most d_i outgoing offers, but can receive any number of acceptable incoming offers. Thus K is an active-demand bound, not a capacity constraint on receiving offers.

Before the next period begins, match lists are cleared. The next period begins with no active relationships.

7. The Outsourcing Decision

7a. Satisfaction Tracking

Each agent i maintains a satisfaction index $s_{i,c}^t$ for each search channel $c \in \{\text{self}, \text{broker}\}$. For the chosen channel, the index updates once per active-demand period as an EWMA of realized match output net of search costs:

$$s_{i,c}^{t+1} = (1 - \omega) s_{i,c}^t + \omega \cdot \tilde{q}$$

where $\omega = 0.2$ and $\tilde{q}$ is the period input in the table. The denominator is requested demand d_i , so unfilled requests contribute zero output. A reciprocal relationship can update two channel-specific states, one for each directed offer.

Channel	Satisfaction input $\tilde{q}$
Self-search	$\frac{\sum q_{ij}}{d_i} - c_s$, summing over accepted directed offers made by i through self-search
Brokered	$\frac{\sum q_{ij} - \phi \cdot n_{i,\text{broker success}}}{d_i}$, summing over accepted directed offers made by i through the broker

Brokered failure gives $\tilde{q} = 0$, self-search failure gives $\tilde{q} = -c_s$.

$s_{i,\text{self}}^t$ covers all self-search, including known neighbors.

At initialization, self-satisfaction is the mean of seed match outcomes and broker-satisfaction is the broker’s seed-data reputation (§7c). Fresh entrants set self-satisfaction to mean neighbor self-satisfaction and broker-satisfaction to current broker reputation.

The first broker choice, regardless of placement success, switches the agent from broker reputation to its own $s_{i,\text{broker}}^t$ (§7b).

7b. Decision Rule

Each agent with $d_i > 0$ compares:

- $\text{score}_{\text{self}} = s_{i,\text{self}}^t$
- $\text{score}_{\text{broker}} = s_{i,\text{broker}}^t$ if the agent has tried the broker, otherwise rep_b^t .

The agent chooses the higher-score channel. Ties are resolved by fair coin flip.

7c. Broker Reputation

$$\text{rep}_b^{t+1} = \begin{cases} \frac{1}{|D^t|} \sum_{i \in D^t} s_{i,\text{broker}}^{t+1} & \text{if } D^t \neq \emptyset \\ \text{rep}_b^t & \text{otherwise} \end{cases}$$

where D^t is the set of agents who outsourced to the broker this period. With current clients, reputation is their mean post-update broker satisfaction; with no clients, it is unchanged. Reputation is initialized from the mean broker seed outcome.

8. Period Measures

Period measures are evaluated after current-period matching, satisfaction, and reputation updates, but before agent entry/exit turnover. Agent decisions do not depend on these measures.

Network Measures

Broker network-position measures use G , including the broker node (§5a). They are computed when $t \bmod \Delta_{\text{net}} = 0$ (default 20). Other period measures are recorded each period.

Betweenness centrality. Freeman betweenness (Freeman, 1977) for the broker node:

$$\text{BC}_b = \frac{1}{(n-1)(n-2)} \sum_{u \neq b} \sum_{v \neq u} \frac{\sigma_{uv}(b)}{\sigma_{uv}}$$

where $n = N + 1$, σ_{uv} is the number of shortest paths from u to v , and $\sigma_{uv}(b)$ is the number passing through b . The double sum counts undirected pairs from both directions, so the denominator is $(n-1)(n-2)$ (Brandes, 2001, p. 9).

Burt’s constraint. Broker ego-network constraint (Burt, 1992):

$$\text{Constraint}_b = \sum_j \left(p_{bj} + \sum_{h \neq b, j} p_{bh} p_{hj} \right)^2$$

where $p_{bj} = 1/d_b$ and $p_{hj} = 1/d_h$ (Everett & Borgatti, 2020). The indirect term uses intermediary h 's degree.

Effective size. Non-redundant broker contacts using the Borgatti (1997) binary-undirected simplification:

$$ES_b = d_b - \frac{2t_b}{d_b},$$

where d_b is the broker's degree and t_b is the number of ties among the broker's neighbors, excluding ties to the broker.

Prediction Quality

Prediction quality is evaluated on random candidate pairs and accepted directed offers:

Sample	Construction	Reported measures
Random candidate pairs	Up to 100 agents with non-empty histories; for each sampled agent i , $\min(40, N - 1)$ non-self partners j . Agent and broker predictions are evaluated against noiseless $Q + f(\mathbf{x}_i, \mathbf{x}_j)$. Both models use the same sample, which is deterministic conditional on seed and period.	Per-agent R^2 , RMSE, bias, and Spearman rank correlation, averaged across sampled agents.
Selected offers	Accepted directed offers by channel, self-search or brokered. A reciprocal relationship can represent two directed decisions.	Channel-specific R^2 , RMSE, bias, and Spearman rank correlation.

The definitions are

$$R^2 = 1 - \frac{\text{MSE}}{\text{Var}_{\text{pop}}(q)}, \quad \text{RMSE} = \sqrt{\frac{1}{n} \sum (\hat{q} - q)^2}, \quad \text{Bias} = \frac{1}{n} \sum (\hat{q} - q).$$

Rank correlation evaluates the strict ranking produced from predicted output. Candidates are ordered by predicted output, with each set of equal predictions ordered uniformly at random. Agent and broker tie-breaking draws are independent. For random candidate pairs, Spearman's ρ_S compares this produced ranking with expected-output ranks. For selected offers, it compares the produced ranking with realized-output ranks. Equal target values retain their average ranks. Thus a constant predictor produces a random ranking with expected rank correlation zero. In the definitions above, q denotes the relevant target: expected output for random candidate pairs and realized output for selected offers. For selected-offer summaries, R^2 , bias, and rank correlation are undefined when $n < 5$ or $\text{Var}_{\text{pop}}(q) < \sigma_\varepsilon^2/6$, while RMSE is defined whenever $n > 0$. Random-candidate summaries use common valid-agent support for all four measures: an agent-level candidate set that fails the sample-size or variance condition is excluded from the R^2 , RMSE, bias, and rank-correlation means.

Other Measures

Measure	Definition
Access vs. assessment	For each accepted broker-directed offer, record whether receiver j was a direct neighbor of sender i in G at the time of the offer. Non-neighbor receivers indicate access; neighbor receivers indicate assessment.
Realized match output by channel	Average realized match output $\bar{q}_c^t$, where $c \in \{\text{self, brokered}\}$.
Requested-position outcomes	For each channel, record the share of requested positions filled and realized output net of channel costs per requested position. Unfilled positions contribute zero output.
Mean channel satisfaction	Cross-agent means $N^{-1} \sum_i s_{i,\text{self}}^t$ and $N^{-1} \sum_i s_{i,\text{broker}}^t$.
Outsourcing rate	Outsourced requested positions divided by total requested positions.
Broker access size	$\ \mathcal{A}_b^t\ = \ \text{Roster}^t \cup D^t\ $, measured after outsourcing decisions form D^t and before entry/exit.
Counterparty concentration	Median and maximum current-period counterparties per agent.
Agent-node degree summaries	Mean, median, minimum, and maximum degree, excluding the broker node.

Part II. Baseline Values and Reproducibility

9. Baseline Values and Reproducibility

9a. Baseline Parameter Values

The table lists baseline values only. Entries are arranged in two columns; each column follows the document order, with the run length listed first.

The baseline learning function is NN. Ridge and its broker variants are alternate experiments (§3).

The baseline broker provides both access and assessment. The assessment-only and access-only modes are alternate experiments (§6c).

Symbol	Baseline	Defined in	Symbol	Baseline	Defined in
T	500	§1	f_r	0.60	§4a
N	1000	§1	λ_c	0.05	§4c
d	8	§1	k_G	6	§5a
d_γ	8	§1	p_{rewire}	0.1	§5a
σ_x	0.5	§1	η_{exit}	0.02	§5c
Q	2.84	§2a	K	5	§6
ρ	0.50	§2a	p_{demand}	0.50	§6
δ	0.5	§2c	$n_{\text{strangers}}$	10	§6a
σ_ε	0.28	§2a	α_R	0.20	§6b
$\eta_{\text{lr},a}, \eta_{\text{lr},b}$	0.003, 0.03	§3a	p_{roster}	0.02	§6b
$E_{\text{init},a}, E_{\text{init},b}$	100, 100	§3a	ω	0.2	§7a
$E_{\text{steps},a}, E_{\text{steps},b}$	50, 50	§3a	Δ_{net}	20	§8
W_p	40 periods	§3a	M_{train}	2000	§3a

9b. Reproducibility

All model-event randomness flows from a single integer seed. The seed determines type draws, G , matching function parameters ($\mathbf{c}$, $\mathbf{A}$, $\mathbf{B}$), broker seed roster, standing-roster churn and replenishment, random ordering within market score ties, and subsequent model events. Random candidate samples and their independent agent and broker tie-breaking draws are deterministic conditional on seed and period. Simulations are fully reproducible given parameters and seed.

References

- Brandes, U. (2001). A faster algorithm for betweenness centrality. *Journal of Mathematical Sociology*, 25(2), 163–177.
- Borgatti, S. P. (1997). Structural holes: Unpacking Burt’s redundancy measures. *Connections*, 20(1), 35–38.
- Burt, R. S. (1992). *Structural holes: The social structure of competition*. Harvard University Press.
- Everett, M. G., & Borgatti, S. P. (2020). Unpacking Burt’s constraint measure. *Social Networks*, 62, 50–57.
- Freeman, L. C. (1977). A set of measures of centrality based on betweenness. *Sociometry*, 40(1), 35–41.
- Watts, D. J., & Strogatz, S. H. (1998). Collective dynamics of ‘small-world’ networks. *Nature*, 393(6684), 440–442.